

Gaurav Sahoo

Chennai, Tamil Nadu

+91 8667622210, gsahoo2110@gmail.com

<https://github.com/gsahoo211004>

[linkedin.com/in/gaurav-sahoo-64221826b](https://www.linkedin.com/in/gaurav-sahoo-64221826b)

Education

Vellore Institute of Technology, Vellore

Sep 2022 – Jul 2027

M.Tech (Integrated) in Computer Science – Data Science Specialization

CGPA: 8.54

Modern Senior Secondary School, Chennai (CBSE)

Apr 2008 – Jul 2022

10th std – 91%

12th std – 92.5%

Experience

• **Research Intern, CDMM, VIT**

May 2025 – Jun 2025

- Built a real-time data analysis platform to track **Random Forest** and **XGBoost** models using **MLflow**.
- Managed the end-to-end pipeline from **data ingestion and experimentation to visualization**.

• **Data Science Intern, Missing Connectz Technology Solutions, Chennai**

Jun 2025 – Jul 2025

- Worked with commercial datasets across data preparation and core **data analysis workflows**.
- Developed automated data pipelines and backend integrations using **Python and Django**.

• **Finance Head, International Society of Automation (ISA), VIT Chapter**

Jan 2024 – Dec 2024

- Managed chapter finances, budgeting, and resource allocation across student-led activities.

Projects

• ***Recovery-Agent – AI Revenue Recovery***

[GitHub](#)

Tech Stack: Python, LangGraph, LightGBM, FastAPI, Pydantic

- Built a LangGraph agent scoring subscription payment-failure risk with a **LightGBM** model (ROC-AUC 0.75).
- Implemented a bounded, compliance-gated escalation playbook (retry → reminder → reauth → discount → escalate) with an append-only audit trail; showed **146× ROI** on a 500-subscriber synthetic batch.

• ***MedScribe-IQ – Clinical Report Auditor & ICD-10 Validator***

[GitHub](#)

Tech Stack: Python, spaCy, Gemini LLM, MLflow, Evidently AI, Streamlit, Docker

- Built a clinical text pipeline using **spaCy**, custom medical dictionaries, and negation detection.
- Added **ICD-10 validation** via the WHO API with a local fallback; integrated **Gemini LLM** for audit summaries.
- Tracked experiments with **MLflow**, monitored data drift with **Evidently AI**, and deployed via **Streamlit + Docker**.

• ***SkillMatch – AI Resume vs Job Fit Evaluator***

[GitHub](#)

Tech Stack: Python, PyMuPDF, SBERT, spaCy, Gemini API, Streamlit, Docker

- Developed a resume evaluator using **PyMuPDF** for PDF parsing and Gemini API for job-role analysis.
- Implemented **SBERT (all-MiniLM-L6-v2)** semantic matching and **spaCy**-based skill extraction with structured skill-gap recommendations.
- Deployed with **Streamlit** and containerized using **Docker**.

• ***EcoDash – Air Quality & Pollution Dashboard***

Tech Stack: Power BI

- Built an interactive **Power BI** dashboard visualizing air quality and pollution trends.

Technical Skills

• **Languages:** Python, SQL, Java, C/C++, R, JavaScript

• **Core Concepts:** Data Structures & Algorithms, OS, OOP, DBMS, Computer Networks, Software Engineering

• **Data Science & ML:** Scikit-learn, PyTorch (Hugging Face/SBERT), Statistical Modelling, Deep Learning

• **NLP & GenAI:** spaCy, NLP Pipelines, NER, LLM/GenAI API Integration, Prompt Engineering, RAG, LangChain

• **MLOps & Data Engineering:** MLflow, Evidently AI, Docker, GitHub Actions CI/CD, ETL/ELT, LangGraph

• **Tools & Platforms:** Git/GitHub, Streamlit, PostgreSQL, MySQL, MongoDB, Linux (Ubuntu), VS Code, Jupyter, Google Colab, Power BI, Excel

Achievements & Certifications

• **Certification:** Machine Learning A-Z: AI, Python & R [2025] — Udemy, Kirill Eremenko & Team

• **Languages:** English, Hindi, Odia, Tamil