

Gaurav Sahoo

gsahoo2110@gmail.com – +91 8667622210 – [GitHub](#) – [LinkedIn](#)

Education

Vellore Institute of Technology, Vellore

Integrated M.Tech in Computer Science – Data Science Specialization

CGPA: **8.54**

Sep 2022 – Jul 2027

Modern Senior Secondary School, Chennai (CBSE)

10th – 91% (CBSE)

12th – 92.5% (CBSE)

Apr 2008 – Jul 2022

Technical Skills

- **Languages:** Python, SQL, Java, C/C++, R, JavaScript
- **Core Concepts:** Data Structures & Algorithms, OS, OOP, DBMS, Computer Networks, Software Engineering
- **Data Science & ML:** Scikit-learn, PyTorch (Hugging Face/SBERT), Statistical Modelling, Deep Learning
- **NLP & GenAI:** spaCy, NLP Pipelines, Named Entity Recognition, LLM/GenAI API Integration, Prompt Engineering, RAG, LangChain
- **MLOps & Data Engineering:** MLflow, Evidently AI, Docker, GitHub Actions CI/CD, ETL/ELT, LangGraph
- **Tools & Platforms:** Git/GitHub, Streamlit, PostgreSQL, MySQL, MongoDB, Linux (Ubuntu), VS Code, Jupyter, Google Colab, Power BI

Projects

• *Recovery-Agent – AI Revenue Recovery*

[GitHub](#)

Tech Stack: Python, LangGraph, LightGBM, FastAPI, Pydantic

- Built a LangGraph agent that scores subscription payment-failure risk using a **LightGBM** model (ROC-AUC 0.75).
- Implemented a bounded, compliance-gated escalation playbook (retry → reminder → reauth → discount → escalate).
- Added an append-only audit trail; demonstrated **146× ROI** across a 500-subscriber batch.

• *MedScribe-IQ – Clinical Report Auditor & ICD-10 Validator*

[GitHub](#)

Tech Stack: Python, spaCy, Gemini LLM, MLflow, Evidently AI, Streamlit, Docker

- Built a clinical text pipeline using **spaCy**, custom medical dictionaries, and negation detection.
- Added **ICD-10 validation** through the WHO API with a local fallback mechanism.
- Integrated **Gemini LLM** for audit summaries and tracked experiments with **MLflow**.
- Monitored data drift using **Evidently AI** and deployed the dashboard with **Streamlit + Docker**.

• *SkillMatch – AI Resume vs Job Fit Evaluator*

[GitHub](#)

Tech Stack: Python, PyMuPDF, SBERT, spaCy, Gemini API, Streamlit, Docker

- Developed an AI resume evaluator using **PyMuPDF** for PDF parsing and Gemini API for job-role analysis.
- Implemented **SBERT (all-MiniLM-L6-v2)** for semantic matching between resumes and job descriptions.
- Built **spaCy**-based skill extraction and generated structured skill-gap recommendations.
- Deployed the application using **Streamlit** and containerized it with **Docker**.

Experience

• **Research Intern, CDMM, VIT**

May 2025 - Jun 2025

- Built a real-time data analysis platform to track **Random Forest** and **XGBoost** models using MLflow.
- Managed the end-to-end pipeline from **data ingestion and experimentation to visualization**.

• **Data Science Intern, Missing Connectz Technology Solutions, Chennai**

Jun 2025 - Jul 2025

- Worked with commercial datasets across data preparation and core **data analysis workflows**.
- Developed automated data pipelines and backend integrations using **Django**.

• **Finance Head, International Society of Automation (ISA), VIT Chapter**

Jan 2024 - Dec 2024

- Managed chapter finances, budgeting, and resource allocation across student-led activities.

Achievements & Certification

- **Certification:** Machine Learning A-Z: AI, Python & R [2025] — Udemy, Kirill Eremenko & Team
- **Languages:** English, Hindi, Odia, Tamil